

CNN

October 3, 2024

1 st124092

2 Kaung Htet Cho

3 Multi-Class Classification with Pytorch

Today, we will

- learn how to build and train a model using Pytorch
- learn about MNIST dataset
- experiment with hyper-parameters tuning

Model development Life-cycle: 1. Prepare the data 2. Define the model architecture 3. **Train the model** 4. Evaluate the model 5. Deploy the model

** : what we are going to focus today !

```
[1]: # For our puffer server we need to browse via a proxy!!
import os
# Set HTTP and HTTPS proxy
os.environ['http_proxy'] = 'http://192.41.170.23:3128'
os.environ['https_proxy'] = 'http://192.41.170.23:3128'
```

```
[2]: # !pip3 install tensorboard
# from torch.utils.tensorboard import SummaryWriter
# %load_ext tensorboard
```

```
[3]: import numpy as np
import torch
import matplotlib.pyplot as plt
```

```
[4]: import torchvision
from torchvision import datasets, transforms, models
from torch.utils.data import DataLoader, Subset
```

3.1 Lets download our data

```
[5]: # load the training data
transform = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.
    ↪ Normalize(mean=0.1307, std=0.3081)
])

train_ds = torchvision.datasets.MNIST(root='.', train=True, download=True, ↪
    ↪ transform=transform)
test_ds = torchvision.datasets.MNIST(root='.', train=False, download=True, ↪
    ↪ transform=transform)
```

```
[6]: print(len(train_ds))
      print(len(test_ds))
```

60000

10000

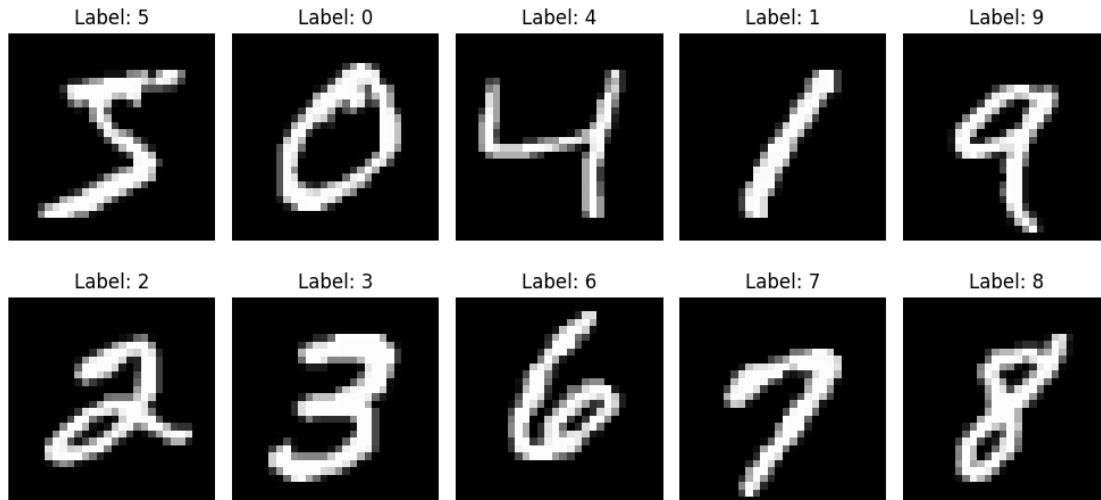
```
[7]: # Create a dictionary to store one image per label
images_per_label = {}

# Loop through the dataset to find one image per label
for img, label in train_ds:
    if label not in images_per_label:
        images_per_label[label] = img
    if len(images_per_label) == 10: # Break the loop once we have all labels
        break

# Plot the images, one per label
fig, axes = plt.subplots(2, 5, figsize=(10, 5))

for i, (label, img) in enumerate(images_per_label.items()):
    ax = axes[i // 5, i % 5]
    ax.imshow(img.squeeze(), cmap='gray')
    ax.set_title(f'Label: {label}')
    ax.axis('off')

plt.tight_layout()
plt.show()
```



```
[8]: # Hyperparameters
lr = 0.01
batch_size = 64
num_epoch = 10
classes = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
```

```
[9]: train_ds = list(train_ds)[:10000]
train_loader = torch.utils.data.DataLoader(train_ds,
                                             batch_size=batch_size,
                                             shuffle=True,
                                             num_workers=2)

test_loader = torch.utils.data.DataLoader(test_ds,
                                           batch_size=batch_size,
                                           shuffle=False,
                                           num_workers=2)
```

```
[10]: train_ds[0][0].shape
```

```
[10]: torch.Size([1, 28, 28])
```

4 1. Define the model

5 Convolutional Neural Network (CNN)

```
[11]: import torch.nn as nn
import torch.nn.functional as F
```

```
[12]: device = 'cuda:0' if torch.cuda.is_available() else 'cpu'
      print(device)
```

cuda:0

The model architecture that we are going to build

Input => conv1 => maxpooling => FC => output

```
[13]: class MyCNN(nn.Module):          # define own MyCNN which inherits nn.Module
      ↪ as a base class
      def __init__(self):
          super(MyCNN, self).__init__()
          self.conv1 = nn.Conv2d(1, 16, kernel_size=3) # conv2d(in_channel,
      ↪ out_channel, kernel_size, stride=1, padding=0, bias=True)
          # (28+2*0-3)/1 + 1 = 26
          self.maxpool = nn.MaxPool2d(2)
          # (26/2)
          self.fc1 = nn.Linear(13*13*16, 10) # Flattened output from convolution
      ↪ followed by pooling layer

      def forward(self, x):
          x = F.relu(self.conv1(x))
          x = self.maxpool(x)
          x = torch.flatten(x,1) # feature maps flattened to 1D tensor output. second
      ↪ dimension (1) refers to flattening across the features, leaving the batch
      ↪ dimension intact.
          x = self.fc1(x)          # produces an output of size 10 (one score for each
      ↪ class)
          return x, F.log_softmax(x, dim=1) # raw output from x (logits), log of
      ↪ softmax of x which normalizes the logits into prob.
```

```
[14]: cnn_model = MyCNN()
      cnn_model = cnn_model.to(device)
      optimizer = torch.optim.SGD(cnn_model.parameters(), lr=lr)
      loss_fn = nn.CrossEntropyLoss()
```

```
[15]: print(cnn_model)
```

```
MyCNN(
  (conv1): Conv2d(1, 16, kernel_size=(3, 3), stride=(1, 1))
  (maxpool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceiling_mode=False)
  (fc1): Linear(in_features=2704, out_features=10, bias=True)
)
```

```

[16]: train_losses = []
      train_accuracies = []
      test_losses = []
      test_accuracies = []

      def train():
          cnn_model.train() # sets model to
          ↪ training mode (dropout and batchnorm behaves in this mode)
          train_corr, train_total, train_running_loss = 0, 0, 0 # counters for
          ↪ tracking training accuracy, total examples, running loss

          for step, (data, y) in enumerate(train_loader): # loops over batch
          ↪ of data in train_loader
              data, y = data.to(device), y.to(device)
              optimizer.zero_grad() # resets gradients
          ↪ to prevent accumulation
              _, logits = cnn_model(data) # gets the logits
              loss = loss_fn(logits, y) # calculates loss
          ↪ comparing with true label
              loss.backward() # back propagation
          ↪ is performed to compute gradients
              optimizer.step() # optimizer
          ↪ updates model params

              y_pred = torch.argmax(logits, 1) # selects the
          ↪ predicted class (the index with the highest value)
              train_corr += torch.sum(torch.eq(y_pred, y).float()).item() #
          ↪ counts correct predictions
              train_total += len(data) # tracks total no.
          ↪ of samples
              train_running_loss += loss.item() # accumulates loss

          # Calculate average loss and accuracy for this epoch
          epoch_loss = train_running_loss / len(train_loader)
          epoch_accuracy = train_corr / train_total

          # Append to lists for plotting
          train_losses.append(epoch_loss)
          train_accuracies.append(epoch_accuracy)

          print(f'Epoch [{epoch+1}] Train Loss: {epoch_loss:.4f}, Accuracy:
          ↪ {epoch_accuracy:.4f}')

      #####

      def test():

```

```

cnn_model.eval() # sets model to
↪evaluation mode
test_corr, test_total, test_running_loss = 0, 0, 0
with torch.no_grad():
    for step, (data, y) in enumerate(test_loader):
        data, y = data.to(device), y.to(device)
        _, logits = cnn_model(data)
        loss = loss_fn(logits, y)
        y_pred = torch.argmax(logits, 1)
        test_corr += torch.sum(torch.eq(y_pred, y).float()).item()
        test_total += len(data)
        test_running_loss += loss.item()
    # Calculate average loss and accuracy for this epoch
    epoch_loss = test_running_loss / len(test_loader)
    epoch_accuracy = test_corr / test_total

    # Append to lists for plotting
    test_losses.append(epoch_loss)
    test accuracies.append(epoch_accuracy)

    print(f'Epoch [{epoch+1}] Test Loss: {epoch_loss:.4f}, Accuracy:
↪{epoch_accuracy:.4f}')

```

```

[17]: for epoch in range(num_epoch):
    print(f"----- Train EPOCH {epoch} -----")
    train()
    print(f"----- Test EPOCH {epoch} -----")
    test()

```

```

----- Train EPOCH 0 -----
Epoch [1] Train Loss: 0.6997, Accuracy: 0.8257
----- Test EPOCH 0 -----
Epoch [1] Test Loss: 0.4318, Accuracy: 0.8756
----- Train EPOCH 1 -----
Epoch [2] Train Loss: 0.3448, Accuracy: 0.9017
----- Test EPOCH 1 -----
Epoch [2] Test Loss: 0.3562, Accuracy: 0.8941
----- Train EPOCH 2 -----
Epoch [3] Train Loss: 0.2968, Accuracy: 0.9146
----- Test EPOCH 2 -----
Epoch [3] Test Loss: 0.3320, Accuracy: 0.9017
----- Train EPOCH 3 -----
Epoch [4] Train Loss: 0.2701, Accuracy: 0.9239
----- Test EPOCH 3 -----
Epoch [4] Test Loss: 0.3043, Accuracy: 0.9107
----- Train EPOCH 4 -----
Epoch [5] Train Loss: 0.2476, Accuracy: 0.9314
----- Test EPOCH 4 -----

```

```


Epoch [5] Test Loss: 0.2672, Accuracy: 0.9219
----- Train EPOCH 5 -----
Epoch [6] Train Loss: 0.2315, Accuracy: 0.9342
----- Test EPOCH 5 -----
Epoch [6] Test Loss: 0.2448, Accuracy: 0.9297
----- Train EPOCH 6 -----
Epoch [7] Train Loss: 0.2167, Accuracy: 0.9381
----- Test EPOCH 6 -----
Epoch [7] Test Loss: 0.2662, Accuracy: 0.9239
----- Train EPOCH 7 -----
Epoch [8] Train Loss: 0.2032, Accuracy: 0.9418
----- Test EPOCH 7 -----
Epoch [8] Test Loss: 0.2345, Accuracy: 0.9317
----- Train EPOCH 8 -----
Epoch [9] Train Loss: 0.1905, Accuracy: 0.9454
----- Test EPOCH 8 -----
Epoch [9] Test Loss: 0.2652, Accuracy: 0.9225
----- Train EPOCH 9 -----
Epoch [10] Train Loss: 0.1803, Accuracy: 0.9500
----- Test EPOCH 9 -----
Epoch [10] Test Loss: 0.2155, Accuracy: 0.9397

```


```

[18]: # Plot the training and test loss
plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Train Loss')
plt.plot(test_losses, label='Test Loss')
plt.title('Loss Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

```



```
[19]: # Plot the training and test accuracy
plt.figure(figsize=(10, 5))
plt.plot(train_accuracies, label='Train Accuracy')
plt.plot(test_accuracies, label='Test Accuracy')
plt.title('Accuracy Curve')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



5.0.1 Plot some prediction

```
[20]: cnn_model.eval()
data, y = next(iter(test_loader))

# 1. push the data to the selected device
data, y = data.to(device), y.to(device)

# 2. feed the data into the model and the model makes predictions
_, logits = cnn_model(data) # raw prediction before applying softmax ;
    ↳ unnormalised scores for each class

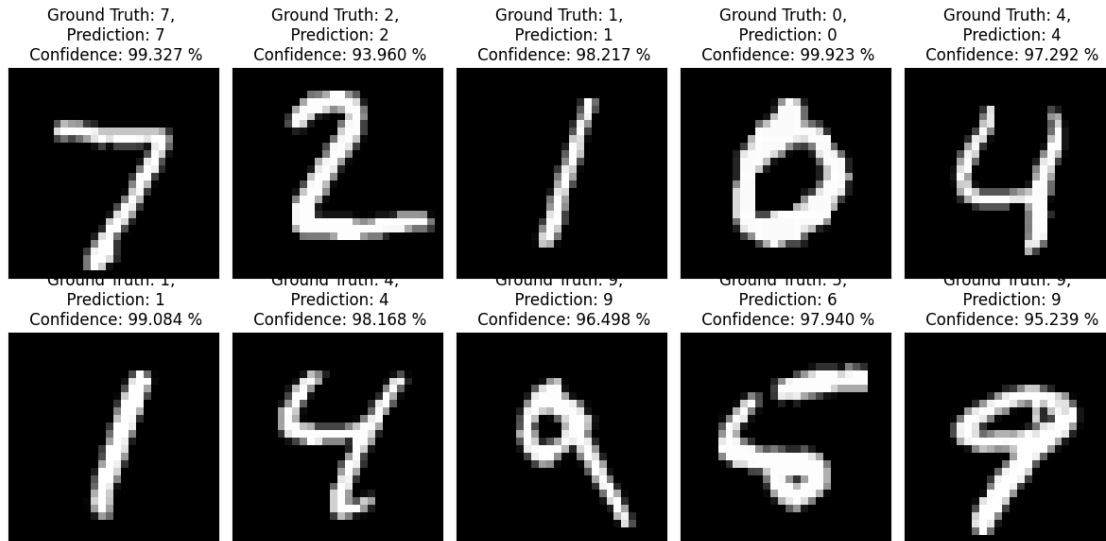
# 3. get the class with highest prob.
y_pred = torch.argmax(logits, 1) # finds the index of the class with the
    ↳ highest value (i.e., the predicted class) along dimension 1,

get_prob = torch.nn.Softmax(dim=1) # converts logits into probabilities that
    ↳ sums to 1
prob = get_prob(logits) # prob is a tensor where each row
    ↳ corresponds to a sample, and each column contains the probability of that
    ↳ sample belonging to a particular class.

# Plot
fig = plt.figure(figsize=(12,6))
for i in range(10):
    plt.subplot(2, 5, i+1)
    plt.imshow(data[i].cpu().detach().numpy().reshape((28,28)), cmap='gray')

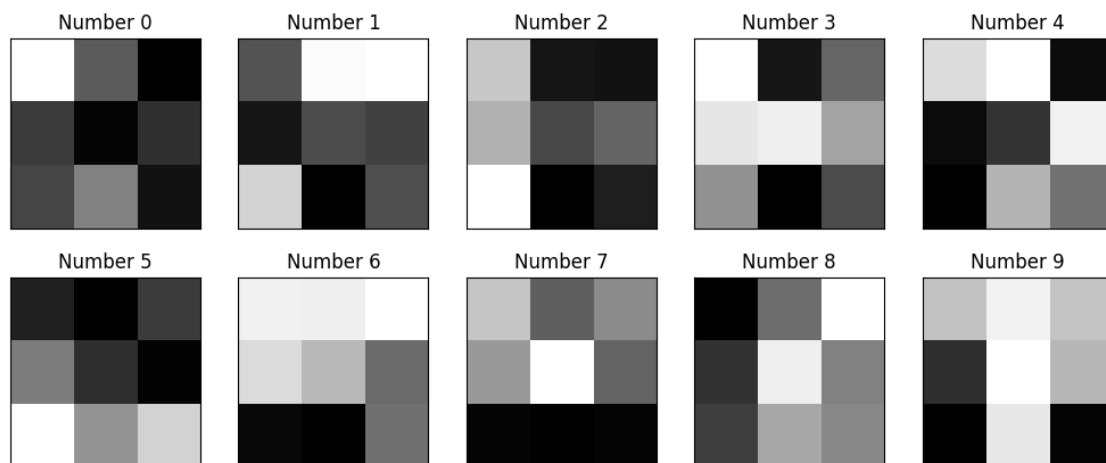
    # detach(): Detaches the tensor from the computation graph, so no gradients
    ↳ are tracked.
    # cpu(): Moves the tensor back to the CPU (important if you're using a GPU).
    # numpy(): Converts the tensor to a NumPy array.

    plt.title(f"Ground Truth: {y[i].cpu().detach().numpy()}, \n Prediction:
    ↳ {y_pred[i].cpu().detach().numpy()} \n Confidence: {prob[i][y_pred[i]] * 100:.
    ↳ 3F} %")
    plt.xticks([])
    plt.yticks([])
plt.tight_layout() # Adjusts the subplot parameters to make sure that subplots
    ↳ fit into the figure area nicely, avoiding overlaps.
plt.show()
```



5.0.2 Visualize filter weights

```
[21]: fig = plt.figure(figsize=[5*2.5, 2*2.5])
for i in range(10):
    # loops through first 10 filters,
    # Each filter corresponds to a learned weight matrix, which is applied to the
    # input image.
    ax = fig.add_subplot(2, 5, i+1)
    # access the ith weight, reshapes to 3*3 matrix
    ws = cnn_model.conv1.weight[i].reshape([3, 3]).cpu().detach().numpy() #
    # Change here if your filter size is changed
    ax.imshow(ws, cmap='gray')
    plt.title(f"Number {i}")
    plt.xticks([])
    plt.yticks([])
```



5.0.3 Visualize feature map


```
[22]: # Visualize feature maps
activation = {} # initializes empty dictionary to store the feature maps

# returns a hook function,
def get_activation(name):
    # hook will capture the layer's output (output) and store it in the
    ↪ activation dictionary with the specified name.
    def hook(model, input, output):
        activation[name] = output.detach()
    return hook

cnn_model.conv1.register_forward_hook(get_activation('conv1')) # This line
    ↪ registers the hook on the first convolutional layer (conv1).
data, _ = test_ds[1] # retrieves a
    ↪ single test sample from the dataset test_ds (data,label)
data.unsqueeze_(0) # Since this is a single image, unsqueeze_(0) changes
    ↪ its shape from [1, 28, 28] to [1, 1, 28, 28], where 1 is the batch size
output = cnn_model(data.to(device)) # output is not required for this case
    ↪ since we stored activations

fm_cov1 = activation['conv1'].squeeze().cpu().detach().numpy() # .squeeze():
    ↪ Removes the extra batch dimension added earlier, so the feature map has the
    ↪ shape [out_channels, height, width].
fig = plt.figure(figsize=[5*2.5, 2*2.5])
for i in range(10):
    ax = fig.add_subplot(2, 5, i+1)
    ax.imshow(fm_cov1[i], cmap='gray')

# The feature maps are the result of applying the learned filters to the input
    ↪ image, so they represent specific patterns or structures detected by the
    ↪ filters.
```



5.1 Lets now try using RESNET18 model

- We can load a pretrained model on Imagenet dataset or train from scratch. For now we are using a pretrained model.

5.1.1 Keypoints:

1. Customizing the Final Layer: Since ResNet-18's final fully connected (fc) layer is designed for ImageNet (1000 classes), we modify it to suit our dataset by setting `resnet18.fc = nn.Linear(resnet18.fc.in_features, num_classes)`.
2. Transformations: The input image size for ResNet-18 is 224x224, so we resize the CIFAR-10 images (originally 32x32) using `transforms.Resize(224)`.
3. Training and Testing: The model is trained using `train_model` and evaluated using `test_model`.

```
[23]: # Load ResNet-18 pre-trained model
from torchvision.models import ResNet18_Weights
resnet18 = models.resnet18(weights=ResNet18_Weights.IMAGENET1K_V1) # or use .
    ↳ DEFAULT for the latest weights

# This code will raise the deprecation warning:
# model = models.resnet18(pretrained=True)
```

```
[24]: # Modify the final layer to match the number of classes (for example, CIFAR-10
    ↳ has 10 classes)
num_classes = 10
resnet18.fc = nn.Linear(resnet18.fc.in_features, num_classes)
```

```
[25]: # Transfer the model to the GPU if available
resnet18 = resnet18.to(device)
```

```
# Define transforms (resize to 224x224 since ResNet expects that input size)
transform = transforms.Compose([
    transforms.Resize(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
])
```

```
[26]: # Download CIFAR-10 dataset (or use your own dataset)
train_dataset = datasets.CIFAR10(root='./data', train=True,
    ↪transform=transform, download=True)
test_dataset = datasets.CIFAR10(root='./data', train=False,
    ↪transform=transform, download=True)

# Subsample the training and test datasets

# Function to subsample CIFAR-10 dataset
def subsample_dataset(dataset, sample_size=1000):
    indices = np.random.choice(len(dataset), sample_size, replace=False)
    subset = Subset(dataset, indices)
    return subset

sample_size = 1000
train_subset = subsample_dataset(train_dataset, sample_size=sample_size)
test_subset = subsample_dataset(train_dataset, sample_size=int(sample_size * 0.
    ↪4))

# Load the data
train_loader = torch.utils.data.DataLoader(dataset=train_subset, batch_size=64,
    ↪shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_subset, batch_size=64,
    ↪shuffle=False)
```

Files already downloaded and verified

Files already downloaded and verified

```
[27]: print(resnet18)
```

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
    bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
    track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1,
    ceil_mode=False)
  (layer1): Sequential(
```

```

        (0): BasicBlock(
          (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
          (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
          (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
        (1): BasicBlock(
          (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
          (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
          (relu): ReLU(inplace=True)
          (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
          (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (layer2): Sequential(
      (0): BasicBlock(
        (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
          (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
      (1): BasicBlock(
        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)

```

```

        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  )
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)

```

```

    )
)
(1): BasicBlock(
  (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
  (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
)
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in_features=512, out_features=10, bias=True)
)

```

```

[28]: # Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(resnet18.parameters(), lr=0.001)

# Training function
def train_model(model, train_loader, criterion, optimizer, num_epochs=5):
    model.train()
    for epoch in range(num_epochs):
        running_loss = 0.0
        for images, labels in train_loader:
            images, labels = images.to(device), labels.to(device)

            # Forward pass
            outputs = model(images)
            loss = criterion(outputs, labels)

            # Backward pass and optimization
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            running_loss += loss.item()

        print(f"Epoch [{epoch+1}/{num_epochs}], Loss: {running_loss/len(train_loader):.4f}")

```

```

[29]: def test_model(model, test_loader):
    model.eval()
    correct = 0

```



```

total = 0
with torch.no_grad():
    for images, labels in test_loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs.data, 1)
        total += labels.size(0)
        correct += (predicted == labels).sum().item()

print(f'Accuracy of the model on the test images: {100 * correct / total:.
↪2f}%')
```

```

[30]: # Training the model
train_model(resnet18, train_loader, criterion, optimizer, num_epochs=5)
```

```

Epoch [1/5], Loss: 1.2798
Epoch [2/5], Loss: 0.5322
Epoch [3/5], Loss: 0.2812
Epoch [4/5], Loss: 0.1641
Epoch [5/5], Loss: 0.0950
```

```

[31]: # Testing the trained model
test_model(resnet18, test_loader)
```

Accuracy of the model on the test images: 73.50%

5.2 Fine-tuning vs. Feature Extraction

- **Fine-tuning:** During fine-tuning, you update the weights of **all layers** in the network during training. This is typically done when you want to adapt a pre-trained model to a new task. By default, when calling `optimizer.step()` on all parameters, the weights of all layers are updated.
- **Feature Extraction:** In feature extraction, you freeze the weights of the pre-trained layers and only train the final layer (or a few newly added layers). This allows the model to use the learned features from the pre-trained network while adjusting the output to the new task.

To freeze the layers in PyTorch, you can set the `requires_grad` attribute of the parameters to `False` like this:

```
“python for param in resnet18.parameters(): param.requires_grad = False
```

6 Take Home Exercise

7 Assignment: Custom CNN on MNIST Dataset

7.1 Objective

In this assignment, you will implement and train a custom Convolutional Neural Network (CNN) to classify handwritten digits from the MNIST dataset.

7.2 Task Breakdown

7.2.1 Part 1: Data Preparation

1. Load and Visualize Data:

- Load the MNIST dataset using PyTorch's `torchvision.datasets` and visualize some sample images along with their labels.
- **Deliverable:** Submit a grid of at least 25 sample images and their respective labels from the MNIST dataset.

7.2.2 Part 2: Custom CNN Implementation

1. Define Your CNN Model:

- Implement a custom CNN class using PyTorch's `torch.nn.Module`. The network should consist of:
 - 2 convolutional layers
 - 2 max-pooling layers
 - At least 1 fully connected layer
 - Use ReLU activations and appropriate dropout layers.

Hints:

- First conv layer: Input channels = (grayscale image), Output channels = 16, Kernel size = 3x3.
- Second conv layer: Output channels = 32, Kernel size = 3x3.
- Max-pooling layers with 2x2 window.

Deliverable: Submit your `MyCNN` class code.

7.2.3 Part 3: Model Training and Evaluation

1. Training the Model:

- Train the CNN on the MNIST training set.
- Use Cross-Entropy Loss and an optimizer (e.g., SGD or Adam).
- Plot the training and validation loss curves over 10-20 epochs.

Deliverable: Submit code for training and a plot showing the loss curves.

2. Accuracy Evaluation:

- Evaluate the model on the MNIST test set and report the accuracy.

Deliverable: Submit the final accuracy on the test set.

7.2.4 Part 4: Visualization

1. Feature Map Visualization:

- Register a forward hook on the second convolutional layer to capture the feature maps.

- Plot the feature maps for a few random test images (at least 3 different digits).

Deliverable: Submit code and images showing feature maps from the second convolutional layer.

7.3 Submission Requirements

- Submit a Jupyter Notebook with the following sections:
 1. Data Loading and Visualization
 2. CNN Model Implementation
 3. Training and Loss Curves
 4. Final Test Accuracy
 5. Feature Map Visualizations

7.4 Evaluation Criteria

- Correctness and clarity of the CNN implementation.
- Proper visualizations of data, loss and accuracy curves, and feature maps.

7.4.1 Part 1: Data Preparation

1. Load and Visualize Data:

- Load the MNIST dataset using PyTorch's `torchvision.datasets` and visualize some sample images along with their labels.
- **Deliverable:** Submit a grid of at least 25 sample images and their respective labels from the MNIST dataset.

```
[32]: import numpy as np
import torch
import matplotlib.pyplot as plt
```

```
[33]: import torchvision
from torchvision import datasets, transforms, models
from torch.utils.data import DataLoader, Subset

# load the training data
transform = torchvision.transforms.Compose([
    torchvision.transforms.ToTensor(),
    torchvision.transforms.
    ↪ Normalize(mean=0.1307, std=0.3081)
])

train_ds = torchvision.datasets.MNIST(root='.', train=True, download=True,
    ↪ transform=transform)
test_ds = torchvision.datasets.MNIST(root='.', train=False, download=True,
    ↪ transform=transform)
```

```
[63]: # Create a list to store 25 random images and their labels
random_images = []
random_labels = []

# Loop through the dataset to find 25 random images
for img, label in train_ds:
    random_images.append(img)
    random_labels.append(label)
    if len(random_images) == 25: # Break the loop once we have 25 images
        break


# Plot the 25 images
fig, axes = plt.subplots(5, 5, figsize=(10, 10))
print(f'type_random_images: {type(random_images[0])}')
print(f'shape_random_images: {random_images[0].shape}')

print(f'shape_random_images_squeeze: {random_images[0].squeeze().shape}')

for i in range(25):
    ax = axes[i // 5, i % 5]
    ax.imshow(random_images[i].squeeze(), cmap='gray')
    ax.set_title(f'Label: {random_labels[i]}')
    ax.axis('off')

plt.tight_layout()
plt.show()
```

```
type_random_images: <class 'torch.Tensor'>
shape_random_images: torch.Size([1, 28, 28])
shape_random_images_squeeze: torch.Size([28, 28])
```



7.4.2 Part 2: Custom CNN Implementation

1. Define Your CNN Model:

- Implement a custom CNN class using PyTorch's `torch.nn.Module`. The network should consist of:
 - 2 convolutional layers
 - 2 max-pooling layers
 - At least 1 fully connected layer
 - Use ReLU activations and appropriate dropout layers.

Hints:

- First conv layer: Input channels = (grayscale image), Output channels = 16, Kernel size = 3x3.
- Second conv layer: Output channels = 32, Kernel size = 3x3.
- Max-pooling layers with 2x2 window.

Deliverable: Submit your MyCNN class code.

```
[35]: import torch
import torch.nn as nn
import torch.nn.functional as F
```

```
[36]: device = 'cuda' if torch.cuda.is_available() else 'cpu'
device
```

```
[36]: 'cuda'
```

```
[37]: class MyCNN(nn.Module):
    def __init__(self):
        super(MyCNN, self).__init__()
        # First Convolution Block
        self.conv1 = nn.Conv2d(1, 16, kernel_size=3) # 28x28x1 -> 26x26x16
        self.dropout1 = nn.Dropout2d(0.25)

        # Second Convolution Block
        self.conv2 = nn.Conv2d(16, 32, kernel_size=3) # 13x13x16 -> 11x11x32
        self.dropout2 = nn.Dropout2d(0.25)

        # Fully Connected Layers
        self.fc1 = nn.Linear(5*5*32, 128) # Calculated from final conv output
        self.dropout3 = nn.Dropout(0.5)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        # First Convolution Block
        x = self.conv1(x) # Convolution
        x = F.relu(x) # ReLU activation
        x = F.max_pool2d(x, 2) # Max pooling
        x = self.dropout1(x) # Dropout

        # Second Convolution Block
        x = self.conv2(x) # Convolution
        x = F.relu(x) # ReLU activation
        x = F.max_pool2d(x, 2) # Max pooling
        x = self.dropout2(x) # Dropout

        # Flatten
        x = torch.flatten(x, 1) # Flatten all dimensions except batch
```

```

    # Fully Connected Layers
    x = self.fc1(x)           # First FC layer
    x = F.relu(x)            # ReLU activation
    x = self.dropout3(x)     # Dropout
    x = self.fc2(x)          # Output layer

    return x, F.log_softmax(x, dim=1)

```

7.4.3 Part 3: Model Training and Evaluation

1. Training the Model:

- Train the CNN on the MNIST training set.
- Use Cross-Entropy Loss and an optimizer (e.g., SGD or Adam).
- Plot the training and validation loss curves over 10-20 epochs.

Deliverable: Submit code for training and a plot showing the loss curves.

2. Accuracy Evaluation:

- Evaluate the model on the MNIST test set and report the accuracy.

Deliverable: Submit the final accuracy on the test set.

```

[38]: cnn_model = MyCNN()
      cnn_model = cnn_model.to(device)

      optimizer = torch.optim.SGD(cnn_model.parameters(), lr=lr)
      loss_fn = nn.CrossEntropyLoss()

```

```

[39]: cnn_model

```

```

[39]: MyCNN(
  (conv1): Conv2d(1, 16, kernel_size=(3, 3), stride=(1, 1))
  (dropout1): Dropout2d(p=0.25, inplace=False)
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
  (dropout2): Dropout2d(p=0.25, inplace=False)
  (fc1): Linear(in_features=800, out_features=128, bias=True)
  (dropout3): Dropout(p=0.5, inplace=False)
  (fc2): Linear(in_features=128, out_features=10, bias=True)
)

```

```

[40]: # Hyperparameters
      lr = 0.01
      batch_size = 64
      num_epoch = 20
      classes = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']

```

```

[41]: train_losses = []
      train_accuracies = []

```

```

test_losses = []
test_accuracies = []

def train():
    cnn_model.train() # sets model to
    ↪training mode (dropout and batchnorm behaves in this mode)
    train_corr, train_total, train_running_loss = 0, 0, 0 # counters for
    ↪tracking training accuracy, total examples, running loss

    for step, (data, y) in enumerate(train_loader): # loops over batch
    ↪of data in train_loader
        data, y = data.to(device), y.to(device)
        optimizer.zero_grad() # resets gradients
    ↪to prevent accumulation
        _, logits = cnn_model(data) # gets the logits
        loss = loss_fn(logits, y) # calculates loss
    ↪comparing with true label
        loss.backward() # back propagation
    ↪is performed to compute gradients
        optimizer.step() # optimizer
    ↪updates model params

        y_pred = torch.argmax(logits, 1) # selects the
    ↪predicted class (the index with the highest value)
        train_corr += torch.sum(torch.eq(y_pred, y).float()).item() #
    ↪counts correct predictions
        train_total += len(data) # tracks total no.
    ↪of samples
        train_running_loss += loss.item() # accumulates loss

    # Calculate average loss and accuracy for this epoch
    epoch_loss = train_running_loss / len(train_loader)
    epoch_accuracy = train_corr / train_total

    # Append to lists for plotting
    train_losses.append(epoch_loss)
    train_accuracies.append(epoch_accuracy)

    print(f'Epoch [{epoch+1}] Train Loss: {epoch_loss:.4f}, Accuracy:
    ↪{epoch_accuracy:.4f}')

#####

```

```

[54]: def test(cnn_model, test_loader):
    cnn_model.eval()

```



```

test_corr, test_total = 0, 0

with torch.no_grad(): # Disable gradient calculation
    for data, y in test_loader:
        data, y = data.to(device), y.to(device)

        # cnn_model returns a tuple, and logits is the second element
        output = cnn_model(data)
        if isinstance(output, tuple):
            logits = output[1] # Adjust this index if necessary
        else:
            logits = output # If it directly returns logits

        y_pred = torch.argmax(logits, 1)
        test_corr += (y_pred == y).sum().item()
        test_total += y.size(0)

accuracy = test_corr / test_total

print(f'Accuracy of the model on the test set: {100 * accuracy:.2f}%')

```

```

[43]: train_ds = list(train_ds)
train_loader = torch.utils.data.DataLoader(train_ds,
                                           batch_size=batch_size,
                                           shuffle=True,
                                           num_workers=2)

test_loader = torch.utils.data.DataLoader(test_ds,
                                           batch_size=batch_size,
                                           shuffle=False,
                                           num_workers=2)

```

```

[44]: for epoch in range(20):
    print(f"----- Train EPOCH {epoch} -----")
    train()

```

```

----- Train EPOCH 0 -----
Epoch [1] Train Loss: 0.8664, Accuracy: 0.7271
----- Train EPOCH 1 -----
Epoch [2] Train Loss: 0.3667, Accuracy: 0.8899
----- Train EPOCH 2 -----
Epoch [3] Train Loss: 0.2722, Accuracy: 0.9207
----- Train EPOCH 3 -----
Epoch [4] Train Loss: 0.2226, Accuracy: 0.9341
----- Train EPOCH 4 -----
Epoch [5] Train Loss: 0.1954, Accuracy: 0.9423
----- Train EPOCH 5 -----
Epoch [6] Train Loss: 0.1753, Accuracy: 0.9483

```

```


----- Train EPOCH 6 -----
Epoch [7] Train Loss: 0.1588, Accuracy: 0.9533
----- Train EPOCH 7 -----
Epoch [8] Train Loss: 0.1490, Accuracy: 0.9559
----- Train EPOCH 8 -----
Epoch [9] Train Loss: 0.1382, Accuracy: 0.9593
----- Train EPOCH 9 -----
Epoch [10] Train Loss: 0.1308, Accuracy: 0.9615
----- Train EPOCH 10 -----
Epoch [11] Train Loss: 0.1231, Accuracy: 0.9636
----- Train EPOCH 11 -----
Epoch [12] Train Loss: 0.1184, Accuracy: 0.9655
----- Train EPOCH 12 -----
Epoch [13] Train Loss: 0.1131, Accuracy: 0.9672
----- Train EPOCH 13 -----
Epoch [14] Train Loss: 0.1089, Accuracy: 0.9688
----- Train EPOCH 14 -----
Epoch [15] Train Loss: 0.1073, Accuracy: 0.9690
----- Train EPOCH 15 -----
Epoch [16] Train Loss: 0.1010, Accuracy: 0.9702
----- Train EPOCH 16 -----
Epoch [17] Train Loss: 0.0973, Accuracy: 0.9717
----- Train EPOCH 17 -----
Epoch [18] Train Loss: 0.0929, Accuracy: 0.9722
----- Train EPOCH 18 -----
Epoch [19] Train Loss: 0.0932, Accuracy: 0.9730
----- Train EPOCH 19 -----
Epoch [20] Train Loss: 0.0912, Accuracy: 0.9737

```


```

[45]: # Plot the training and test loss
plt.figure(figsize=(10, 5))
plt.plot(train_losses, label='Train Loss')
plt.title('Loss Curve')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.show()

```



```
[46]: # Plot the training and test accuracy
plt.figure(figsize=(10, 5))
plt.plot(train_accuracies, label='Train Accuracy')
plt.title('Accuracy Curve')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



```
[56]: test(cnn_model, test_loader)
```

Accuracy of the model on the test set: 98.82%

7.4.4 Part 4: Visualization

1. Feature Map Visualization:

- Register a forward hook on the second convolutional layer to capture the feature maps.
- Plot the feature maps for a few random test images (at least 3 different digits).

Deliverable: Submit code and images showing feature maps from the second convolutional layer.

```
[62]: import torch
import matplotlib.pyplot as plt

activation = {}

def get_activation(name):
    def hook(model, input, output):
        activation[name] = output.detach()
    return hook

cnn_model.conv2.register_forward_hook(get_activation('conv2'))

num_images = 3

for idx in range(num_images):
    data, label = test_ds[idx]    # Get a random test sample
    data.unsqueeze_(0)           # Add batch dimension, so shape becomes [1, 1, 28, 28]
    ↪

    # Forward pass through the network
    output = cnn_model(data.to(device))


    # Get the feature map from conv2 layer
    fm_conv2 = activation['conv2'].squeeze().cpu().detach().numpy() # Shape: 1, 1, 128, 128
    ↪ [num_channels, height, width]

    fig = plt.figure(figsize=(10, 5))


    # Plot the first 10 feature maps
    for i in range(10):
        ax = fig.add_subplot(2, 5, i+1)
        ax.imshow(fm_conv2[i], cmap='gray')
        ax.axis('off')
```

```
plt.suptitle(f"Feature maps from conv2 for test image {idx} (Label:␣  
↪{label})")  
plt.show()
```

Feature maps from conv2 for test image 0 (Label: 7)



Feature maps from conv2 for test image 1 (Label: 2)



Feature maps from conv2 for test image 2 (Label: 1)

