

YOLO

October 27, 2024

```
[26]: #st124092@ait.asia
```

```
[1]: import torch
import torch.nn as nn
```

```
/home/st124092/work/cuda116/.venv/lib/python3.8/site-packages/tqdm/auto.py:21:
TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
```

```
[2]: 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
[2]: 'cuda'
```

1 YOLO ARCHITECTURE

Image Reference: <https://www.datacamp.com/blog/yolo-object-detection-explained>

The following code is inspired and taken from this repo :
https://github.com/aladdinpersson/Machine-Learning-Collection/blob/master/ML/Pytorch/object_detection/YOLOv3.py

1.1 Lets build our model

```
[3]: """
Information about architecture config:
Tuple is structured by (kernel_size, filters, stride, padding)
"M" is simply maxpooling with stride 2x2 and kernel 2x2
List is structured by tuples and lastly int with number of repeats
"""

architecture_config = [
    (7, 64, 2, 3),
    "M",
    (3, 192, 1, 1),
    "M",
    (1, 128, 1, 0),
    (3, 256, 1, 1),
    (1, 256, 1, 0),
```

```

(3, 512, 1, 1),
"M",
[(1, 256, 1, 0), (3, 512, 1, 1), 4], # list architecture
(1, 512, 1, 0),
(3, 1024, 1, 1),
"M",
[(1, 512, 1, 0), (3, 1024, 1, 1), 2],
(3, 1024, 1, 1),
(3, 1024, 2, 1),
(3, 1024, 1, 1),
(3, 1024, 1, 1),
]

```

```

[4]: class CNNBlock(nn.Module):
    def __init__(self, in_channels, out_channels, **kwargs):
        super(CNNBlock, self).__init__()
        self.conv = nn.Conv2d(in_channels, out_channels, bias=False, **kwargs)
        self.batchnorm = nn.BatchNorm2d(out_channels) # Not used in
        original implementation!
        self.leakyrelu = nn.LeakyReLU(0.1)

    def forward(self, x):
        return self.leakyrelu(self.batchnorm(self.conv(x)))

```

```

[5]: class YOLOv1(nn.Module):
    def __init__(self, in_channels=3, **kwargs):
        super(YOLOv1, self).__init__()
        self.architecture = architecture_config
        self.in_channels = in_channels
        self.darknet = self._create_conv_layers(self.architecture)
        self.fcs = self._create_fc(**kwargs)

    def forward(self, x):
        x = self.darknet(x)
        return self.fcs(torch.flatten(x, start_dim=1))

    def _create_conv_layers(self, architecture):
        layers = []
        in_channels = self.in_channels

        for x in architecture:
            if type(x) == tuple:
                layers += [
                    CNNBlock(in_channels, out_channels = x[1],
                        kernel_size=x[0], stride=x[2], padding=x[3],)
                ]
                in_channels = x[1]

```

```

        if type(x) == str:
            layers += [
                nn.MaxPool2d(kernel_size=2, stride=2)]

        if type(x) == list:
            conv1 = x[0]
            conv2 = x[1]
            num_repeats = x[2]

            for _ in range(num_repeats):
                layers += [
                    CNNBlock(in_channels, out_channels = conv1[1],
↪kernel_size= conv1[0], stride=conv1[2], padding=conv1[3],)
                layers += [
                    CNNBlock(conv1[1], out_channels = conv2[1],
↪kernel_size= conv2[0], stride=conv2[2], padding=conv2[3],)
                    in_channels = conv2[1]
            return nn.Sequential(*layers)

    def _create_fc(self, split_size, num_boxes, num_classes):
        S,B,C = split_size, num_boxes, num_classes
        return nn.Sequential(
            nn.Flatten(),
            nn.Linear(1024*S*S , 4096),
            nn.Dropout(0.5),
            nn.LeakyReLU(0.1),
            nn.Linear(4096, S*S*(C + B*5)),
        )

```

[6]: *# Test our model*

```

model = Yolov1(split_size=7, num_boxes=2, num_classes=20)
print(model)

```

```

Yolov1(
  (darknet): Sequential(
    (0): CNNBlock(
      (conv): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3),
bias=False)
      (batchnorm): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (leakyrelu): LeakyReLU(negative_slope=0.1)
    )
    (1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (2): CNNBlock(
      (conv): Conv2d(64, 192, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),

```

```

bias=False)
    (batchnorm): BatchNorm2d(192, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(3): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(4): CNNBlock(
    (conv): Conv2d(192, 128, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batchnorm): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(5): CNNBlock(
    (conv): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (batchnorm): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(6): CNNBlock(
    (conv): Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batchnorm): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(7): CNNBlock(
    (conv): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (batchnorm): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(9): CNNBlock(
    (conv): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
    (batchnorm): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(10): CNNBlock(
    (conv): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
    (batchnorm): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leakyrelu): LeakyReLU(negative_slope=0.1)
)

```

```

(11): CNNBlock(
  (conv): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batchnorm): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(12): CNNBlock(
  (conv): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batchnorm): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(13): CNNBlock(
  (conv): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batchnorm): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(14): CNNBlock(
  (conv): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batchnorm): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(15): CNNBlock(
  (conv): Conv2d(512, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batchnorm): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(16): CNNBlock(
  (conv): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batchnorm): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(17): CNNBlock(
  (conv): Conv2d(512, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
  (batchnorm): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(18): CNNBlock(
  (conv): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)

```

```

        (batchnorm): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leakyrelu): LeakyReLU(negative_slope=0.1)
    )
    (19): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (20): CNNBlock(
        (conv): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batchnorm): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leakyrelu): LeakyReLU(negative_slope=0.1)
    )
    (21): CNNBlock(
        (conv): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batchnorm): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leakyrelu): LeakyReLU(negative_slope=0.1)
    )
    (22): CNNBlock(
        (conv): Conv2d(1024, 512, kernel_size=(1, 1), stride=(1, 1), bias=False)
        (batchnorm): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leakyrelu): LeakyReLU(negative_slope=0.1)
    )
    (23): CNNBlock(
        (conv): Conv2d(512, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batchnorm): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leakyrelu): LeakyReLU(negative_slope=0.1)
    )
    (24): CNNBlock(
        (conv): Conv2d(1024, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
        (batchnorm): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leakyrelu): LeakyReLU(negative_slope=0.1)
    )
    (25): CNNBlock(
        (conv): Conv2d(1024, 1024, kernel_size=(3, 3), stride=(2, 2), padding=(1,
1), bias=False)
        (batchnorm): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (leakyrelu): LeakyReLU(negative_slope=0.1)
    )
    (26): CNNBlock(
        (conv): Conv2d(1024, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1,

```

```

1), bias=False)
    (batchnorm): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (leakyrelu): LeakyReLU(negative_slope=0.1)
)
(27): CNNBlock(
  (conv): Conv2d(1024, 1024, kernel_size=(3, 3), stride=(1, 1), padding=(1,
1), bias=False)
  (batchnorm): BatchNorm2d(1024, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (leakyrelu): LeakyReLU(negative_slope=0.1)
)
)
(fcs): Sequential(
  (0): Flatten(start_dim=1, end_dim=-1)
  (1): Linear(in_features=50176, out_features=4096, bias=True)
  (2): Dropout(p=0.5, inplace=False)
  (3): LeakyReLU(negative_slope=0.1)
  (4): Linear(in_features=4096, out_features=1470, bias=True)
)
)

```

```

[7]: x = torch.randn(2, 3, 448, 448)
      print(model(x).shape)                                # 7*7*30 = 1470

```

```
torch.Size([2, 1470])
```

1.2 Lets build our Loss Function

```

[8]: from dataset import VOCDataset
      from utils import intersection_over_union

```

```

[9]: class YoloLoss(nn.Module):
      """
      Calculate the loss for yolo (v1) model
      """

      def __init__(self, S=7, B=2, C=20):
          super(YoloLoss, self).__init__()
          self.mse = nn.MSELoss(reduction="sum")

          """
          S is split size of image (in paper 7),
          B is number of boxes (in paper 2),
          C is number of classes (in paper and VOC dataset is 20),
          """

          self.S = S
          self.B = B

```

```

self.C = C

# These are from Yolo paper, signifying how much we should
# pay loss for no object (noobj) and the box coordinates (coord)
self.lambda_noobj = 0.5
self.lambda_coord = 5

def forward(self, predictions, target):
    # predictions are shaped (BATCH_SIZE, S*S(C+B*5) when inputted
    predictions = predictions.reshape(-1, self.S, self.S, self.C + self.B * 5)

    # Calculate IoU for the two predicted bounding boxes with target bbox
    iou_b1 = intersection_over_union(predictions[..., 21:25], target[..., 21:25])
    iou_b2 = intersection_over_union(predictions[..., 26:30], target[..., 21:25])
    ious = torch.cat([iou_b1.unsqueeze(0), iou_b2.unsqueeze(0)], dim=0)

    # Take the box with highest IoU out of the two prediction
    # Note that bestbox will be indices of 0, 1 for which bbox was best
    iou_maxes, bestbox = torch.max(ious, dim=0)
    exists_box = target[..., 20].unsqueeze(3) # in paper this is Iobj_i

    # ===== #
    #   FOR BOX COORDINATES   #
    # ===== #

    # Set boxes with no object in them to 0. We only take out one of the
    # predictions, which is the one with highest Iou calculated previously.
    box_predictions = exists_box * (
        (
            bestbox * predictions[..., 26:30]
            + (1 - bestbox) * predictions[..., 21:25]
        )
    )

    box_targets = exists_box * target[..., 21:25]

    # Take sqrt of width, height of boxes to ensure that
    box_predictions[..., 2:4] = torch.sign(box_predictions[..., 2:4]) *
    torch.sqrt(
        torch.abs(box_predictions[..., 2:4] + 1e-6)
    )
    box_targets[..., 2:4] = torch.sqrt(box_targets[..., 2:4])

```



```

box_loss = self.mse(
    torch.flatten(box_predictions, end_dim=-2),
    torch.flatten(box_targets, end_dim=-2),
)

# ===== #
#   FOR OBJECT LOSS   #
# ===== #

# pred_box is the confidence score for the bbox with highest IoU
pred_box = (
    bestbox * predictions[..., 25:26] + (1 - bestbox) * predictions[...
↪, 20:21]
)

object_loss = self.mse(
    torch.flatten(exists_box * pred_box),
    torch.flatten(exists_box * target[..., 20:21])),
)

# ===== #
#   FOR NO OBJECT LOSS   #
# ===== #

#max_no_obj = torch.max(predictions[..., 20:21], predictions[..., 25:
↪26])

#no_object_loss = self.mse(
#    torch.flatten((1 - exists_box) * max_no_obj, start_dim=1),
#    torch.flatten((1 - exists_box) * target[..., 20:21], start_dim=1),
#)

no_object_loss = self.mse(
    torch.flatten((1 - exists_box) * predictions[..., 20:21],
↪start_dim=1),
    torch.flatten((1 - exists_box) * target[..., 20:21], start_dim=1),
)

no_object_loss += self.mse(
    torch.flatten((1 - exists_box) * predictions[..., 25:26],
↪start_dim=1),
    torch.flatten((1 - exists_box) * target[..., 20:21], start_dim=1)
)

# ===== #
#   FOR CLASS LOSS   #
# ===== #

```

```

class_loss = self.mse(
    torch.flatten(exists_box * predictions[..., :20], end_dim=-2,),
    torch.flatten(exists_box * target[..., :20], end_dim=-2,),
)

loss = (
    self.lambda_coord * box_loss # first two rows in paper
    + object_loss # third row in paper
    + self.lambda_noobj * no_object_loss # forth row
    + class_loss # fifth row
)

return loss

```

```

[10]: import torchvision.transforms as transforms
import torch.optim as optim
import torchvision.transforms.functional as FT
from tqdm import tqdm
from torch.utils.data import DataLoader
from utils import (
    non_max_suppression,
    mean_average_precision,
    intersection_over_union,
    cellboxes_to_boxes,
    get_bboxes,
    plot_image,
    save_checkpoint,
    load_checkpoint,
)

```

```

[11]: seed = 123
torch.manual_seed(seed)

```

```

[11]: <torch._C.Generator at 0x7fee807d5610>

```

```

[12]: # Hyperparameters etc.
LEARNING_RATE = 2e-5
DEVICE = "cuda" if torch.cuda.is_available else "cpu"
BATCH_SIZE = 16 # 64 in original paper but I don't have that much vram, grad_
    ↪accum?
WEIGHT_DECAY = 0
EPOCHS = 10
NUM_WORKERS = 2
PIN_MEMORY = True
LOAD_MODEL = False
LOAD_MODEL_FILE = "best_so_far.pth.tar"
IMG_DIR = "dataset/images"

```

```
LABEL_DIR = "dataset/labels"
```

```
[13]: class Compose(object):
        def __init__(self, transforms):
            self.transforms = transforms

        def __call__(self, img, bboxes):
            for t in self.transforms:
                img, bboxes = t(img), bboxes

            return img, bboxes

transform = Compose([transforms.Resize((448, 448)), transforms.ToTensor(),])
```

```
[14]: def train_fn(train_loader, model, optimizer, loss_fn):
        loop = tqdm(train_loader, leave=True)
        mean_loss = []

        for batch_idx, (x, y) in enumerate(loop):
            x, y = x.to(DEVICE), y.to(DEVICE)
            out = model(x)
            loss = loss_fn(out, y)
            mean_loss.append(loss.item())
            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

            # update progress bar
            loop.set_postfix(loss=loss.item())

        print(f"Mean loss was {sum(mean_loss)/len(mean_loss)}")
```

```
[15]: def main():
        model = Yolov1(split_size=7, num_boxes=2, num_classes=20).to(DEVICE)
        optimizer = optim.Adam(
            model.parameters(), lr=LEARNING_RATE, weight_decay=WEIGHT_DECAY
        )
        loss_fn = YoloLoss()

        if LOAD_MODEL:
            load_checkpoint(torch.load(LOAD_MODEL_FILE), model, optimizer)

        train_dataset = VOCDataset(
            "dataset/100examples.csv",
            transform=transform,
            img_dir=IMG_DIR,
```

```

        label_dir=LABEL_DIR,
    )

    test_dataset = VOCDataset(
        "dataset/test.csv", transform=transform, img_dir=IMG_DIR,
↪label_dir=LABEL_DIR,
    )

    train_loader = DataLoader(
        dataset=train_dataset,
        batch_size=BATCH_SIZE,
        num_workers=NUM_WORKERS,
        pin_memory=PIN_MEMORY,
        shuffle=True,
        drop_last=True,
    )

    test_loader = DataLoader(
        dataset=test_dataset,
        batch_size=BATCH_SIZE,
        num_workers=NUM_WORKERS,
        pin_memory=PIN_MEMORY,
        shuffle=True,
        drop_last=True,
    )

    for epoch in range(EPOCHS):
        pred_boxes, target_boxes = get_bboxes(
            train_loader, model, iou_threshold=0.5, threshold=0.4
        )

        mean_avg_prec = mean_average_precision(
            pred_boxes, target_boxes, iou_threshold=0.5, box_format="midpoint"
        )
        print(f"Train mAP: {mean_avg_prec}")

        #if mean_avg_prec > 0.9:
        #    checkpoint = {
        #        "state_dict": model.state_dict(),
        #        "optimizer": optimizer.state_dict(),
        #    }
        #    save_checkpoint(checkpoint, filename=LOAD_MODEL_FILE)
        #    import time
        #    time.sleep(10)

    train_fn(train_loader, model, optimizer, loss_fn)

```

[16]: main()

Train mAP: 0.0

100%| | 6/6 [00:02<00:00, 2.02it/s, loss=715]

Mean loss was 889.3602294921875

Train mAP: 0.0

100%| | 6/6 [00:02<00:00, 2.56it/s, loss=395]

Mean loss was 537.6829884847006

Train mAP: 0.0

100%| | 6/6 [00:02<00:00, 2.43it/s, loss=503]

Mean loss was 478.6702575683594

Train mAP: 0.0

100%| | 6/6 [00:02<00:00, 2.62it/s, loss=526]

Mean loss was 435.08201090494794

Train mAP: 0.0

100%| | 6/6 [00:02<00:00, 2.64it/s, loss=306]

Mean loss was 368.8504587809245

Train mAP: 6.61375597701408e-05

100%| | 6/6 [00:02<00:00, 2.40it/s, loss=281]

Mean loss was 321.43519083658856

Train mAP: 3.561253106454387e-05

100%| | 6/6 [00:02<00:00, 2.28it/s, loss=202]

Mean loss was 282.7339324951172

Train mAP: 0.0

100%| | 6/6 [00:02<00:00, 2.21it/s, loss=231]

Mean loss was 252.1270980834961

Train mAP: 5.733943544328213e-05

100%| | 6/6 [00:02<00:00, 2.44it/s, loss=265]

Mean loss was 251.42701212565103

Train mAP: 0.00028062198543921113

100%| | 6/6 [00:02<00:00, 2.23it/s, loss=214]

Mean loss was 229.84151458740234

1.2.1 write a test function for validation data,

```
[17]: def valid_fn(valid_loader, model, loss_fn):
    model.eval()  # Set model to evaluation mode
    mean_loss = []
    loop = tqdm(valid_loader, leave=True)

    with torch.no_grad():
        for batch_idx, (x, y) in enumerate(loop):
            x, y = x.to(DEVICE), y.to(DEVICE)
            out = model(x)
            loss = loss_fn(out, y)
            mean_loss.append(loss.item())

            # update progress bar
            loop.set_postfix(loss=loss.item())

    print(f"Mean validation loss was {sum(mean_loss)/len(mean_loss)}")
    model.train()
    return sum(mean_loss)/len(mean_loss)
```

1.2.2 train the yolo model with training and validation set on PASCAL VOC data.

```
[18]: from sklearn.model_selection import train_test_split
    from torch.utils.data import Subset
    import torch

    full_dataset = VOCDataset(
        "dataset/train.csv",
        transform=transform,
        img_dir=IMG_DIR,
        label_dir=LABEL_DIR,
```

```

)

dataset_size = len(full_dataset)
indices = list(range(dataset_size))
train_indices, val_indices = train_test_split(
    indices,
    test_size=0.2,
    random_state=seed
)

train_dataset = Subset(full_dataset, train_indices)
val_dataset = Subset(full_dataset, val_indices)

total_size = len(full_dataset)
train_size = len(train_dataset)
val_size = len(val_dataset)

print(f"Total dataset size: {total_size}")
print(f"Training set size: {train_size} ({train_size/total_size*100:.1f}%)")
print(f"Validation set size: {val_size} ({val_size/total_size*100:.1f}%)")

```

```

Total dataset size: 16550
Training set size: 13240 (80.0%)
Validation set size: 3310 (20.0%)

```

```

[19]: def main():
    model = YOLOv1(split_size=7, num_boxes=2, num_classes=20).to(DEVICE)
    optimizer = optim.Adam(
        model.parameters(), lr=LEARNING_RATE, weight_decay=WEIGHT_DECAY
    )
    loss_fn = YOLOLoss()

    if LOAD_MODEL:
        load_checkpoint(torch.load(LOAD_MODEL_FILE), model, optimizer)

    train_loader = DataLoader(
        dataset=train_dataset,
        batch_size=BATCH_SIZE,
        num_workers=NUM_WORKERS,
        pin_memory=PIN_MEMORY,
        shuffle=True,
        drop_last=True,
    )

    val_loader = DataLoader(
        dataset=val_dataset,

```

```

        batch_size=BATCH_SIZE,
        num_workers=NUM_WORKERS,
        pin_memory=PIN_MEMORY,
        shuffle=False,
        drop_last=True,
    )

    for epoch in range(EPOCHS):
        pred_boxes, target_boxes = get_bboxes(
            train_loader, model, iou_threshold=0.5, threshold=0.4
        )

        mean_avg_prec = mean_average_precision(
            pred_boxes, target_boxes, iou_threshold=0.5, box_format="midpoint"
        )
        print(f"Train mAP: {mean_avg_prec}")

        train_fn(train_loader, model, optimizer, loss_fn)
        val_loss = valid_fn(val_loader, model, loss_fn)

        checkpoint = {
            "state_dict": model.state_dict(),
            "optimizer": optimizer.state_dict(),
        }
        save_checkpoint(checkpoint, filename=LOAD_MODEL_FILE)
        print("Model saved after 10 epochs.")

```

```
[20]: !export CUDA_VISIBLE_DEVICES=0,1,2,3
```

```
[21]: main()
```

```

Train mAP: 0.0
100%|      | 827/827 [03:24<00:00,  4.04it/s, loss=192]
Mean loss was 251.90797998135437
100%|      | 206/206 [00:52<00:00,  3.96it/s, loss=163]
Mean validation loss was 163.26945588195207
=> Saving checkpoint
Train mAP: 0.002744792029261589
100%|      | 827/827 [03:02<00:00,  4.53it/s, loss=182]
Mean loss was 179.07935164427383
100%|      | 206/206 [00:44<00:00,  4.64it/s, loss=153]
Mean validation loss was 162.85283397933813
=> Saving checkpoint

```


Train mAP: 0.007169117219746113

100%| | 827/827 [04:43<00:00, 2.92it/s, loss=135]

Mean loss was 161.70346530605255

100%| | 206/206 [01:13<00:00, 2.82it/s, loss=148]

Mean validation loss was 148.96882344218133

=> Saving checkpoint

Train mAP: 0.012085916474461555

100%| | 827/827 [04:05<00:00, 3.37it/s, loss=128]

Mean loss was 152.16088289678314

100%| | 206/206 [01:09<00:00, 2.95it/s, loss=146]

Mean validation loss was 151.80737538013642

=> Saving checkpoint

Train mAP: 0.015234148129820824

100%| | 827/827 [03:06<00:00, 4.43it/s, loss=136]

Mean loss was 145.9815778074991

100%| | 206/206 [02:26<00:00, 1.41it/s, loss=139]

Mean validation loss was 144.69658575706111

=> Saving checkpoint

Train mAP: 0.019742552191019058

100%| | 827/827 [05:48<00:00, 2.38it/s, loss=128]

Mean loss was 142.86766745072612

100%| | 206/206 [00:53<00:00, 3.82it/s, loss=150]

Mean validation loss was 147.16873750408877

=> Saving checkpoint

Train mAP: 0.015162885189056396

100%| | 827/827 [03:35<00:00, 3.83it/s, loss=163]

Mean loss was 136.3023274063055

100%| | 206/206 [00:54<00:00, 3.80it/s, loss=152]

Mean validation loss was 137.13396129793333

=> Saving checkpoint

Train mAP: 0.02364708110690117

100%| | 827/827 [03:16<00:00, 4.20it/s, loss=85.3]

Mean loss was 130.03348006187355

100%| | 206/206 [00:45<00:00, 4.55it/s, loss=130]

```

Mean validation loss was 122.83234498107318
=> Saving checkpoint
Train mAP: 0.03647088259458542

100%|      | 827/827 [03:09<00:00, 4.36it/s, loss=121]

Mean loss was 124.56289040912596

100%|      | 206/206 [00:54<00:00, 3.79it/s, loss=138]

Mean validation loss was 121.81533954212966
=> Saving checkpoint
Train mAP: 0.033590167760849

100%|      | 827/827 [02:51<00:00, 4.83it/s, loss=112]

Mean loss was 118.8215225754799

100%|      | 206/206 [00:50<00:00, 4.10it/s, loss=133]

Mean validation loss was 136.5670263420031
=> Saving checkpoint
Model saved after 10 epochs.

```

1.2.3 the performance of trained model on test data of PASCAL VOC in terms of mAP.

```

[22]: LOAD_MODEL = True
      LOAD_MODEL_FILE = "best_so_far.pth.tar"

[23]: model = Yolov1(split_size=7, num_boxes=2, num_classes=20).to(DEVICE)
      optimizer = optim.Adam(model.parameters(), lr=LEARNING_RATE,
      ↪weight_decay=WEIGHT_DECAY)

      # if LOAD_MODEL:
      #     load_checkpoint(torch.load(LOAD_MODEL_FILE), model, optimizer)
      #     print("Model loaded successfully from", LOAD_MODEL_FILE)

      if LOAD_MODEL:
          checkpoint = torch.load(LOAD_MODEL_FILE)
          load_checkpoint(checkpoint, model, optimizer) # Use the imported function
          print("Model loaded successfully from", LOAD_MODEL_FILE)

      test_dataset = VOCDataset(
          "dataset/test.csv", transform=transform, img_dir=IMG_DIR,
          ↪label_dir=LABEL_DIR,
      )

      test_loader = DataLoader(
          dataset=test_dataset,
          batch_size=BATCH_SIZE,

```

```

num_workers=NUM_WORKERS,
pin_memory=PIN_MEMORY,
shuffle=False,
drop_last=True,
)

```

=> Loading checkpoint
Model loaded successfully from best_so_far.pth.tar

```

[25]: # Define the evaluation function
def evaluate_test_set(test_loader, model, device, iou_threshold=0.5,
    ↪conf_threshold=0.4):
    model.eval()
    pred_boxes, target_boxes = [], []

    model = model.to(device)

    with torch.no_grad():
        pred_boxes, target_boxes = get_bboxes(
            test_loader, model, iou_threshold=iou_threshold,
            ↪threshold=conf_threshold, device=device
        )

    mean_avg_prec = mean_average_precision(
        pred_boxes, target_boxes, iou_threshold=iou_threshold,
        ↪box_format="midpoint"
    )

    print(f"Test Set mAP: {mean_avg_prec:.4f}")

evaluate_test_set(test_loader, model, DEVICE, iou_threshold=0.5,
    ↪conf_threshold=0.4)

```

Test Set mAP: 0.0251